


SmolLM-135M-Instruct w8a8 (I8 + BF16)

483 tensors, 219.8 MB raw → 171.8 MB archived (1.279×). One column per transformer layer, one row per tensor role.


(a) Compression achieved

darker is a smaller archive for that tensor




(b) Program the search selected

colour is the program's outermost operator



(c) Where the bytes go, by tensor role

share of raw bytes removed — entropy coding alone, then what the structural operators add on top



Outside the layer stack: lm_head 56.6 MB 1.51× embed_tokens 56.6 MB 1.51× norm 1.2 kB 1.66×