


TinyLlama-1.1B-Chat (BF16): what Brevis does to each tensor

201 tensors, 2.2 GB raw → 1.5 GB archived (1.510×). One column per transformer layer, one row per tensor role.

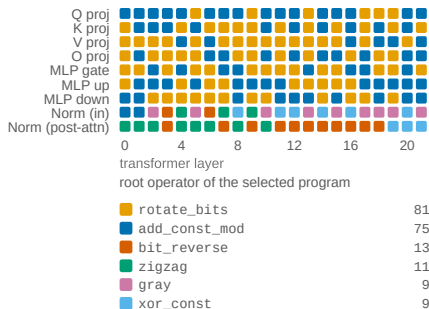
(a) Compression achieved

darker is a smaller archive for that tensor



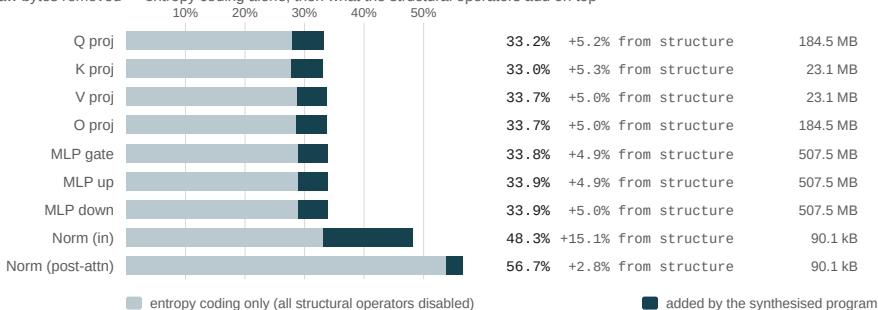
(b) Program the search selected

colour is the program's outermost operator



(c) Where the bytes go, by tensor role

share of raw bytes removed — entropy coding alone, then what the structural operators add on top



Outside the layer stack:

lm_head 131.1 MB 1.51×

embed_tokens 131.1 MB 1.51×

norm 4.1 kB 2.14×