


ViT-Base/16 (F32)

200 tensors, 346.3 MB raw → 288.1 MB archived (1.202×). One column per transformer layer, one row per tensor role.


(a) Compression achieved

darker is a smaller archive for that tensor




(b) Program the search selected

colour is the program's outermost operator



(c) Where the bytes go, by tensor role

share of raw bytes removed — entropy coding alone, then what the structural operators add on top



Outside the layer stack: classifier 3.1 MB 1.20× projection 2.4 MB 1.19× embeddings 605.2 kB 1.18×